

# Capstone Project

## Image classifier for the SVHN dataset

### Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

### How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

### Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: 1 import tensorflow as tf
2 from scipy.io import loadmat
3 from tensorflow.keras.models import Sequential
4 from tensorflow.keras.layers import Conv2D, Flatten, BatchNormalization, MaxPool2D, Dense, Dropout
5 from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping
6 from tensorflow.keras.models import load_model
7
8 import pandas as pd
9 import numpy as np
10 import matplotlib.pyplot as plt
11 from sklearn.model_selection import train_test_split
12 %matplotlib inline
```

SVHN overview image

For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning". NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

```
In [2]: 1 # Run this cell to load the dataset
2
3 train = loadmat('data/train_32x32.mat')
4 test = loadmat('data/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `x` and `y` for the input images and labels respectively.

## 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
In [3]: 1 # Extract features and labels from the dataset
2 X_train_full = train['X']
3 y_train_full = train['y']
4
5 X_test_full = test['X']
6 y_test_full = test['y']
```

```
In [4]: 1 X_train_full.shape, X_test_full.shape
```

```
Out[4]: ((32, 32, 3, 73257), (32, 32, 3, 26032))
```

```
In [5]: 1 # Use a subset of the dataset (10,000 images)
2 subset_size = 10000
3 X_train_full = X_train_full[:, :, :, :subset_size]
4 y_train_full = y_train_full[:subset_size]
5
6 X_test_full = X_test_full[:, :, :, :subset_size]
7 y_test_full = y_test_full[:subset_size]
```

```
In [6]: 1 X_train_full.shape, X_test_full.shape
```

```
Out[6]: ((32, 32, 3, 10000), (32, 32, 3, 10000))
```

```
In [7]: 1 # Display a random sample of original color images before preprocessing
2 fig, axes = plt.subplots(1, 10, figsize=(15, 1.5))
3 for i in range(10):
4     random_index = np.random.randint(0, X_train_full.shape[3])
5     axes[i].imshow(X_train_full[:, :, :, random_index])
6     axes[i].axis('off')
7 plt.show()
```



```
In [8]: 1 # Convert the training and test images to grayscale
2 X_train_gray = np.mean(X_train_full, axis=3, keepdims=True)
3 X_test_gray = np.mean(X_test_full, axis=3, keepdims=True)
```

```
In [9]: 1 # Display a random sample of grayscale images before preprocessing
2 X_train_gray = np.mean(X_train_full, axis=3, keepdims=True)
3
4 fig, ax = plt.subplots(1, 10, figsize=(15, 1.5))
5 for i in range(10):
6     ax[i].set_axis_off()
7     ax[i].imshow(np.squeeze(X_train_gray[i]), cmap="gray")
8
9 plt.show()
```



```
In [10]: 1 # Preprocess the data
2 X_train_full = X_train_full / 255.0 # Normalize pixel values to be between 0 and 1
3 X_test_full = X_test_full / 255.0

In [11]: 1 # Reshape the data to have 3 channels (for RGB) and add channel dimension
2 X_train_full = X_train_full.reshape((X_train_full.shape[3], 32, 32, 3)).transpose(0, 2, 1, 3)
3 X_test_full = X_test_full.reshape((X_test_full.shape[3], 32, 32, 3)).transpose(0, 2, 1, 3)

In [12]: 1 # Convert labels to one-hot encoding
2 y_train_full[y_train_full == 10] = 0 # Replace label 10 with 0
3 y_train_full = tf.keras.utils.to_categorical(y_train_full, 10)

In [13]: 1 y_test_full[y_test_full == 10] = 0 # Replace label 10 with 0
2 y_test_full = tf.keras.utils.to_categorical(y_test_full, 10)

In [14]: 1 # Split the data into training and validation sets
2 X_train, X_val, y_train, y_val = train_test_split(X_train_full, y_train_full, test_size=0.2, random_state=42)

In [15]: 1 X_train.shape, y_train.shape
Out[15]: ((8000, 32, 32, 3), (8000, 10))
```

## 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [16]: 1 # Build the MLP model
2 mlp_model = Sequential([
3     Flatten(input_shape=(32, 32, 3)),
4     Dense(256, activation='relu'),
5     Dense(128, activation='relu'),
6     Dense(64, activation='relu'),
7     Dense(10, activation='softmax')
8 ])
```

```
In [17]: 1 # Print the model summary
2 mlp_model.summary()

Model: "sequential"
```

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 3072)	0
dense (Dense)	(None, 256)	786688
dense_1 (Dense)	(None, 128)	32896
dense_2 (Dense)	(None, 64)	8256
dense_3 (Dense)	(None, 10)	650
Total params: 828,490		
Trainable params: 828,490		
Non-trainable params: 0		

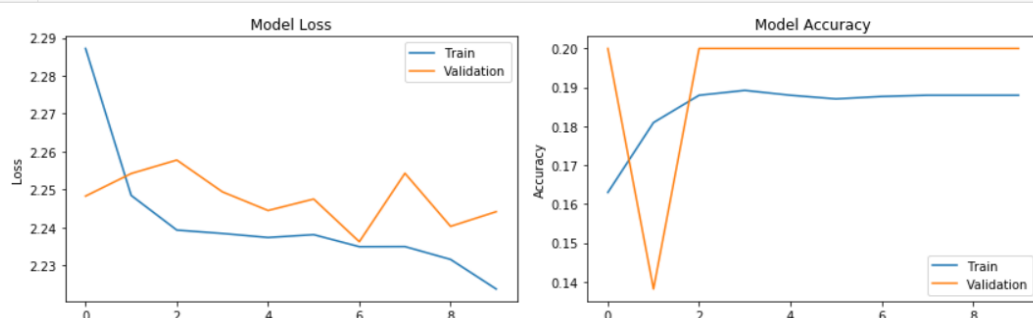
```
In [18]: 1 # Compile the model
2 mlp_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [19]: 1 # Set up callbacks
2 checkpoint_callback = ModelCheckpoint("best_mlp_model.h5", save_best_only=True)
3 early_stopping_callback = EarlyStopping(patience=5, restore_best_weights=True)
```

```
In [20]: 1 # Train the model
2 history = mlp_model.fit(X_train, y_train, epochs=10, validation_split=0.2, callbacks=[checkpoint_callback, early_stopping_callback])
```

```
Train on 6400 samples, validate on 1600 samples
Epoch 1/10
6400/6400 [=====] - 11s 2ms/sample - loss: 2.2872 - accuracy: 0.1630 - val_loss: 2.2483 - val_accuracy: 0.2000
Epoch 2/10
6400/6400 [=====] - 8s 1ms/sample - loss: 2.2485 - accuracy: 0.1809 - val_loss: 2.2543 - val_accuracy: 0.1381
Epoch 3/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2393 - accuracy: 0.1880 - val_loss: 2.2578 - val_accuracy: 0.2000
Epoch 4/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2385 - accuracy: 0.1892 - val_loss: 2.2494 - val_accuracy: 0.2000
Epoch 5/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2374 - accuracy: 0.1880 - val_loss: 2.2445 - val_accuracy: 0.2000
Epoch 6/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2382 - accuracy: 0.1870 - val_loss: 2.2475 - val_accuracy: 0.2000
Epoch 7/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2350 - accuracy: 0.1877 - val_loss: 2.2363 - val_accuracy: 0.2000
Epoch 8/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2350 - accuracy: 0.1880 - val_loss: 2.2543 - val_accuracy: 0.2000
Epoch 9/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2316 - accuracy: 0.1880 - val_loss: 2.2403 - val_accuracy: 0.2000
Epoch 10/10
6400/6400 [=====] - 9s 1ms/sample - loss: 2.2238 - accuracy: 0.1880 - val_loss: 2.2442 - val_accuracy: 0.2000
```

```
In [21]: 1 # Plot learning curves
2 plt.figure(figsize=(12, 4))
3
4 # Plot training & validation loss values
5 plt.subplot(1, 2, 1)
6 plt.plot(history.history['loss'])
7 plt.plot(history.history['val_loss'])
8 plt.title('Model Loss')
9 plt.xlabel('Epoch')
10 plt.ylabel('Loss')
11 plt.legend(['Train', 'Validation'], loc='upper right')
12
13 # Plot training & validation accuracy values
14 plt.subplot(1, 2, 2)
15 plt.plot(history.history['accuracy'])
16 plt.plot(history.history['val_accuracy'])
17 plt.title('Model Accuracy')
18 plt.xlabel('Epoch')
19 plt.ylabel('Accuracy')
20 plt.legend(['Train', 'Validation'], loc='lower right')
21
22 plt.tight_layout()
23 plt.show()
```



### 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.)*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [22]: 1 cnn_model = Sequential([
2         # Convolutional Block 1
3         Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)),
4         BatchNormalization(),
5         MaxPool2D((2, 2)),
6         Dropout(0.25),
7
8         # Convolutional Block 2
9         Conv2D(64, (3, 3), activation='relu'),
10        BatchNormalization(),
11        MaxPool2D((2, 2)),
12        Dropout(0.25),
13
14        # Convolutional Block 3
15        Conv2D(128, (3, 3), activation='relu'),
16        BatchNormalization(),
17        MaxPool2D((2, 2)),
18        Dropout(0.25),
19
20        # Flatten and Dense Layers
21        Flatten(),
22        Dense(512, activation='relu'),
23        BatchNormalization(),
24        Dropout(0.5),
25
26        # Output Layer
27        Dense(10, activation='softmax')
28    ])
```


```
In [23]: 1 cnn_model.summary()

Model: "sequential_1"

Layer (type)                Output Shape              Param #
=====
conv2d (Conv2D)              (None, 30, 30, 32)       896
batch_normalization (Batch Normalization) (None, 30, 30, 32)       128
max_pooling2d (MaxPooling2D) (None, 15, 15, 32)       0
dropout (Dropout)            (None, 15, 15, 32)       0
conv2d_1 (Conv2D)            (None, 13, 13, 64)       18496
batch_normalization_1 (Batch Normalization) (None, 13, 13, 64)       256
max_pooling2d_1 (MaxPooling2D) (None, 6, 6, 64)        0
dropout_1 (Dropout)          (None, 6, 6, 64)        0
conv2d_2 (Conv2D)            (None, 4, 4, 128)       73856
batch_normalization_2 (Batch Normalization) (None, 4, 4, 128)       512
max_pooling2d_2 (MaxPooling2D) (None, 2, 2, 128)       0
dropout_2 (Dropout)          (None, 2, 2, 128)       0
flatten_1 (Flatten)          (None, 512)              0
dense_4 (Dense)              (None, 512)              262656
batch_normalization_3 (Batch Normalization) (None, 512)              2048
dropout_3 (Dropout)          (None, 512)              0
dense_5 (Dense)              (None, 10)               5130
=====
Total params: 363,978
Trainable params: 362,506
Non-trainable params: 1,472
```

```
In [24]: 1 # Compile the model
2 cnn_model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
In [25]: 1 # Model Checkpoint and Early Stopping
2 early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)
```

In [27]: 



Train on 8000 samples, validate on 2000 samples

In [29]: In [31]: [illegible]

Test Loss: 2.2892  
Test Accuracy: 14.41%







```
In [50]: 1 # Display the images with true labels and predicted distributions for CNN
2 plt.figure(figsize=(15, 8))
3
4 for i in range(5):
5     plt.subplot(2, 5, i + 1)
6     plt.imshow(selected_images[i])
7     plt.title(f"True Label: {class_names[selected_labels[i]]}")
8     plt.axis('off')
9
10    plt.subplot(2, 5, i + 6)
11    plt.bar(range(10), cnn_predictions[i], color='orange', alpha=0.7)
12    plt.xticks(range(10), class_names)
13    plt.xlabel('Class')
14    plt.ylabel('Probability')
15    plt.title(f"CNN Prediction")
16
17 plt.tight_layout()
18 plt.show()
```

