

Capstone Project

January 1, 2024

1 Capstone Project

1.1 Probabilistic generative models

1.1.1 Instructions

In this notebook, you will practice working with generative models, using both normalising flow networks and the variational autoencoder algorithm. You will create a synthetic dataset with a normalising flow with randomised parameters. This dataset will then be used to train a variational autoencoder, and you will use the trained model to interpolate between the generated images. You will use concepts from throughout this course, including Distribution objects, probabilistic layers, bijectors, ELBO optimisation and KL divergence regularisers.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

1.1.2 How to submit

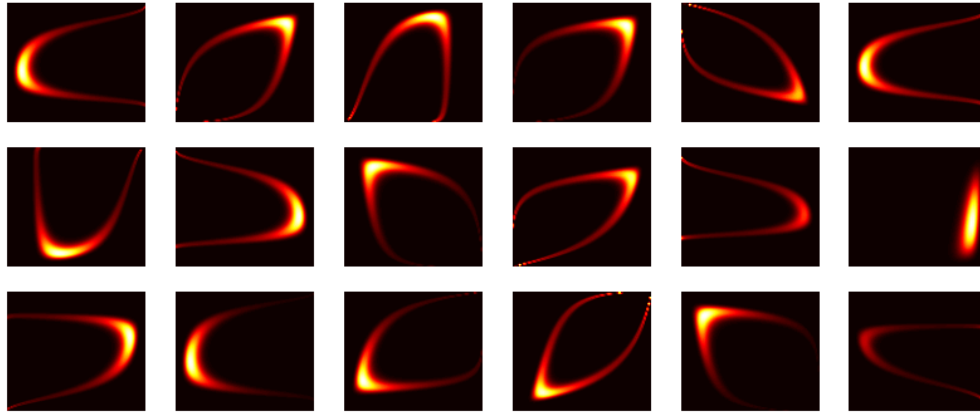
When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

1.1.3 Let's get started!

We'll start by running some imports below. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: import tensorflow as tf
import tensorflow_probability as tfp
tfd = tfp.distributions
tfb = tfp.bijectors
tfpl = tfp.layers

import numpy as np
import matplotlib.pyplot as plt
```



Flags overview image

```
%matplotlib inline

from tqdm import tqdm

from tensorflow.keras.layers import InputLayer, Conv2D, Flatten, Dense, Reshape, BatchNormalization
from tensorflow.keras.models import Sequential, Model
```

For the capstone project, you will create your own image dataset from contour plots of a transformed distribution using a random normalising flow network. You will then use the variational autoencoder algorithm to train generative and inference networks, and synthesise new images by interpolating in the latent space.

The normalising flow

- To construct the image dataset, you will build a normalising flow to transform the 2-D Gaussian random variable $z = (z_1, z_2)$, which has mean $\mathbf{0}$ and covariance matrix $\Sigma = \sigma^2 \mathbf{I}_2$, with $\sigma = 0.3$.
- This normalising flow uses bijectors that are parameterised by the following random variables:
 - $\theta \sim U[0, 2\pi)$
 - $a \sim N(3, 1)$

The complete normalising flow is given by the following chain of transformations: * $f_1(z) = (z_1, z_2 - 2)$, * $f_2(z) = (z_1, \frac{z_2}{2})$, * $f_3(z) = (z_1, z_2 + az_1^2)$, * $f_4(z) = Rz$, where R is a rotation matrix with angle θ , * $f_5(z) = \tanh(z)$, where the \tanh function is applied elementwise.

The transformed random variable x is given by $x = f_5(f_4(f_3(f_2(f_1(z)))))$. * You should use or construct bijectors for each of the transformations f_i , $i = 1, \dots, 5$, and use `tfb.Chain` and `tfb.TransformedDistribution` to construct the final transformed distribution. * Ensure to implement the `log_det_jacobian` methods for any subclassed bijectors that you write. * Display a scatter plot of samples from the base distribution. * Display 4 scatter plot images of the transformed distribution from your random normalising flow, using samples of θ and a . Fix the axes of these 4 plots to the range $[-1, 1]$.


```
In [2]: # base distribution
```

```
# To construct the image dataset, you will build a normalising flow to transform the 2
# which has mean 0 and covariance matrix = 2I2, with = 0.3.
mu, sigma = 0, 0.3
base_dist = tfd.MultivariateNormalDiag(loc=[mu, mu], scale_diag=[sigma, sigma])

# This normalising flow uses bijectors that are parameterised by the following random
# * U[0,2)
theta_dist = tfd.Uniform(low=0, high=2*np.pi)
# * aN(3,1)
a_dist = tfd.Normal(loc=3, scale=1)
```

```
In [3]: # Display a scatter plot of samples from the base distribution.
```

```
dist_plot = base_dist.sample(300).numpy().squeeze()
plt.figure()
plt.scatter(dist_plot[:, 0], dist_plot[:, 1])
plt.title("Base Distribution")
plt.show()
```



```
In [4]: # polynomial bijector (f3)
```

```
class Polynomial(tfb.Bijector):
    def __init__(self, a, name="Polynomial", validate_args=False):
```

```

        super(Polynomial, self).__init__(validate_args=validate_args,
                                          forward_min_event_ndims=1,
                                          is_constant_jacobian=True,
                                          name=name)

        self.a = tf.cast(a, dtype=tf.float32)

    def _forward(self, x):
        x = tf.cast(x, dtype=tf.float32)
        return tf.concat([x[..., 0:1],
                          x[..., 1:] + self.a * tf.square(x[..., 0:1])], axis=-1)

    def _inverse(self, y):
        y = tf.cast(y, dtype=tf.float32)
        return tf.concat([y[..., 0:1],
                          y[..., 1:] - self.a * tf.square(y[..., 0:1])], axis=-1)

# Ensure to implement the log_det_jacobian methods for any subclassed bijectors th
    def _forward_log_det_jacobian(self, x):
        return tf.constant(0., dtype=x.dtype)

```

In [5]: *# rotation bijector (f4)*

```

class Rotation(tfb.Bijector):
    def __init__(self, theta, validate_args=False, name="Rotation"):
        super(Rotation, self).__init__(validate_args=validate_args,
                                       forward_min_event_ndims=1,
                                       name=name)

        self.rot_matrix = tf.convert_to_tensor([[tf.cos(theta), -tf.sin(theta)],
                                                [tf.sin(theta), tf.cos(theta)]], dtype=

    def _forward(self, x):
        x = tf.cast(x, dtype=tf.float32)
        return tf.linalg.matvec(self.rot_matrix, x)

    def _inverse(self, y):
        y = tf.cast(y, dtype=tf.float32)
        return tf.linalg.matvec(tf.transpose(self.rot_matrix), y)

# Ensure to implement the log_det_jacobian methods for any subclassed bijectors th
    def _forward_log_det_jacobian(self, x):
        return tf.constant(0., dtype=x.dtype)

```

In [6]: *# chained bijectors*

```

def GetFlow(theta, a):
    # The complete normalising flow is given by the following chain of transformations

```

```

# *  $f_1(z)=(z_1, z_2)$ ,
f1 = tfb.Shift([0, -2])
# *  $f_2(z)=(z_1, z_2)$ ,
f2 = tfb.Scale([1, 0.5])
# *  $f_3(z)=(z_1, z_2+az_1)$ ,
f3 = Polynomial(a)
# *  $f_4(z)=Rz$ , where  $R$  is a rotation matrix with angle ,
f4 = Rotation(theta)
# *  $f_5(z)=\tanh(z)$ , where the tanh function is applied elementwise.
f5 = tfb.Tanh()

# The transformed random variable  $x$  is given by  $x=f_5(f_4(f_3(f_2(f_1(z)))))$ .
return tfb.Chain([f5, f4, f3, f2, f1])

# You should use or construct bijectors for each of the transformations  $f_i$ ,  $i=1, \dots, 5$ ,
# and use tfb.Chain and tfb.TransformedDistribution to construct the final transformed
FlowDist = lambda theta, a, base_dist: tfd.TransformedDistribution(distribution=base_d
                                                                    bijector=GetFlow


```

In [7]: # Display 4 scatter plot images of the transformed distribution from your random normal
Fix the axes of these 4 plots to the range `[1,1]`.

```

plt.figure(figsize=(10, 10))
for i in range(4):
    theta = theta_dist.sample().numpy()
    a = a_dist.sample().numpy()
    flow_dist = FlowDist(theta, a, base_dist)
    plt.subplot(2, 2, i+1)
    samples = flow_dist.sample(300).numpy().squeeze()
    plt.scatter(samples[:,0], samples[:, 1])
    plt.title("theta = {:.2f}, a = {:.2f}".format(theta, a))
    plt.xlim([-1,1])
    plt.ylim([-1,1])
plt.show()

```



1.2 2. Create the image dataset

- You should now use your random normalising flow to generate an image dataset of contour plots from your random normalising flow network.
 - Feel free to get creative and experiment with different architectures to produce different sets of images!
- First, display a sample of 4 contour plot images from your normalising flow network using 4 independently sampled sets of parameters.
 - You may find the following `get_densities` function useful: this calculates density values for a (batched) Distribution for use in a contour plot.
- Your dataset should consist of at least 1000 images, stored in a numpy array of shape $(N, 36, 36, 3)$. Each image in the dataset should correspond to a contour plot of a transformed dis-

tribution from a normalising flow with an independently sampled set of parameters s, T, S, b . It will take a few minutes to create the dataset.

- As well as the `get_densities` function, the `get_image_array_from_density_values` function will help you to generate the dataset.
 - This function creates a numpy array for an image of the contour plot for a given set of density values Z . Feel free to choose your own options for the contour plots.
- Display a sample of 20 images from your generated dataset in a figure.

In [8]: *# Helper function to compute transformed distribution densities*

```
X, Y = np.meshgrid(np.linspace(-1, 1, 100), np.linspace(-1, 1, 100))
inputs = np.transpose(np.stack((X, Y)), [1, 2, 0])
```

```
def get_densities(transformed_distribution):
    """
    This function takes a (batched) Distribution object as an argument, and returns a
    numpy array Z of shape (batch_shape, 100, 100) of density values, that can be used to make a
    contour plot with:
    plt.contourf(X, Y, Z[b, ...], cmap='hot', levels=100)
    where b is an index into the batch shape.
    """
    batch_shape = transformed_distribution.batch_shape
    Z = transformed_distribution.prob(np.expand_dims(inputs, 2))
    Z = np.transpose(Z, list(range(2, 2+len(batch_shape))) + [0, 1])
    return Z
```

In [9]: *# Helper function to convert contour plots to numpy arrays*

```
import numpy as np
from matplotlib.backends.backend_agg import FigureCanvasAgg as FigureCanvas
from matplotlib.figure import Figure

def get_image_array_from_density_values(Z):
    """
    This function takes a numpy array Z of density values of shape (100, 100)
    and returns an integer numpy array of shape (36, 36, 3) of pixel values for an image.
    """
    assert Z.shape == (100, 100)
    fig = Figure(figsize=(0.5, 0.5))
    canvas = FigureCanvas(fig)
    ax = fig.gca()
    ax.contourf(X, Y, Z, cmap='hot', levels=100)
    ax.axis('off')
    fig.tight_layout(pad=0)

    ax.margins(0)
    fig.canvas.draw()
```

```

image_from_plot = np.frombuffer(fig.canvas.tostring_rgb(), dtype=np.uint8)
image_from_plot = image_from_plot.reshape(fig.canvas.get_width_height()[::-1] + (3,))
return image_from_plot


```

In [10]: *# First, display a sample of 4 contour plot images from your normalising flow network*

```

plt.figure(figsize=(20,5))
for i in range(4):
    theta = theta_dist.sample().numpy()
    a = a_dist.sample().numpy()
    flow_dist = FlowDist(theta, a, base_dist)
    flow_dist = tfd.BatchReshape(flow_dist, [1])
    plt.subplot(1, 4, i+1)
    # You may find the following get_densities function useful: this calculates densities
    plt.contourf(X, Y, get_densities(flow_dist).squeeze(), cmap='hot', levels=50)
    plt.title("theta = {:.2f}, a = {:.2f}".format(theta, a))
    plt.xlim([-1,1])
    plt.ylim([-1,1])
plt.tight_layout()
plt.show()

```



In [11]: *# Your dataset should consist of at least 1000 images, stored in a numpy array of shape (1000, 100, 100, 3)*
Each image in the dataset should correspond to a contour plot of a transformed distribution

```

images = []
img_params = []

for _ in tqdm(range(1000)):
    theta = theta_dist.sample().numpy()
    a = a_dist.sample().numpy()
    flow_dist = FlowDist(theta, a, base_dist)
    flow_dist = tfd.BatchReshape(flow_dist, [1])
    # As well as the get_densities function, the get_image_array_from_density_values function is useful
    densities = get_densities(flow_dist).squeeze()
    images.append(get_image_array_from_density_values(densities))


images = np.array(images)

```


100%|| 1000/1000 [07:04<00:00, 2.36it/s]

In [12]: *# Display a sample of 20 images from your generated dataset in a figure.*

```
plt.figure(figsize=(20,5))
for i in range(20):
    plt.subplot(2, 10, i+1)
    plt.imshow(images[i])
    plt.axis("off")
plt.tight_layout()
plt.show()
```



1.3 3. Make `tf.data.Dataset` objects

- You should now split your dataset to create `tf.data.Dataset` objects for training and validation data.
- Using the `map` method, normalise the pixel values so that they lie between 0 and 1.
- These Datasets will be used to train a variational autoencoder (VAE). Use the `map` method to return a tuple of input and output Tensors where the image is duplicated as both input and output.
- Randomly shuffle the training Dataset.
- Batch both datasets with a batch size of 20, setting `drop_remainder=True`.
- Print the `element_spec` property for one of the Dataset objects.

In [13]: *# You should now split your dataset to create tf.data.Dataset objects for training and*

```
dataset = tf.data.Dataset.from_tensor_slices(images)
split_size = int(len(images)*0.75)
print(split_size)

train_set = dataset.take(split_size)
val_set = dataset.skip(split_size)
```

750

```

In [14]: def prepare_data(dataset: tf.data.Dataset):

    # Using the map method, normalise the pixel values so that they lie between 0 and 1
    dataset = dataset.map(lambda x: tf.cast(x, tf.float32))
    dataset = dataset.map(lambda x: x/255.0)
    # These Datasets will be used to train a variational autoencoder (VAE). Use the map
    dataset = dataset.map(lambda x: (x,x))

    return dataset

train_set = prepare_data(train_set)
# Randomly shuffle the training Dataset.
train_set = train_set.shuffle(split_size)
val_set = prepare_data(val_set)

# Batch both datasets with a batch size of 20, setting drop_remainder=True.
train_set = train_set.batch(batch_size=20, drop_remainder=True)
val_set = val_set.batch(batch_size=20, drop_remainder=True)

In [15]: # Print the element_spec property for one of the Dataset objects.
train_set.element_spec

Out [15]: (TensorSpec(shape=(20, 36, 36, 3), dtype=tf.float32, name=None),
          TensorSpec(shape=(20, 36, 36, 3), dtype=tf.float32, name=None))

```

1.4 4. Build the encoder and decoder networks

- You should now create the encoder and decoder for the variational autoencoder algorithm.
- You should design these networks yourself, subject to the following constraints:
 - The encoder and decoder networks should be built using the Sequential class.
 - The encoder and decoder networks should use probabilistic layers where necessary to represent distributions.
 - The prior distribution should be a zero-mean, isotropic Gaussian (identity covariance matrix).
 - The encoder network should add the KL divergence loss to the model.
- Print the model summary for the encoder and decoder networks.

```

In [16]: # You should now create the encoder and decoder for the variational autoencoder algorithm
# You should design these networks yourself, subject to the following constraints:
# * The encoder and decoder networks should be built using the Sequential class.
# * The encoder and decoder networks should use probabilistic layers where necessary

latent_dim = 2
image_dim = images.shape[1:]

# * The prior distribution should be a zero-mean, isotropic Gaussian (identity covariance)
prior = tfd.MultivariateNormalDiag(loc=tf.Variable(tf.zeros(latent_dim), dtype=tf.float32), dtype=tf.float32)

```

```
scale_diag = tfp.util.TransformedVariable(initial_
bijector
```

```
In [17]: # encoder part
# inspired by https://www.tensorflow.org/tutorials/generative/cvae
encoder = Sequential([
    InputLayer(input_shape=image_dim),

    Conv2D(filters=32, kernel_size=(3,3), activation='relu', padding=
    BatchNormalization(),
    Conv2D(filters=64, kernel_size=(3,3), activation='relu', padding=
    BatchNormalization(),
    Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding=
    BatchNormalization(),
    Conv2D(filters=8, kernel_size=(1,1), activation='relu', padding=
    BatchNormalization(),

    Flatten(),
    Dense(100),
    BatchNormalization(),

    Dense(tfppl.MultivariateNormalTriL.params_size(latent_dim), acti
    tfppl.MultivariateNormalTriL(latent_dim),
    # * The encoder network should add the KL divergence loss to th
    tfppl.KLDivergenceAddLoss(prior,
                                use_exact_kl = False,
                                test_points_fn = lambda q:q.sample(5),
                                test_points_reduce_axis=(0,1))

], name='encoder')
# Print the model summary for the encoder and decoder networks.
encoder.summary()
```

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/linear_ops.py:114: tf.nn.conv2d is deprecated and will be removed in a future version. Instructions for updating:
Do not pass `graph_parents`. They will no longer be used.
Model: "encoder"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 36, 36, 32)	896
batch_normalization (Batch Normalization)	(None, 36, 36, 32)	128
conv2d_1 (Conv2D)	(None, 36, 36, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 36, 36, 64)	256
conv2d_2 (Conv2D)	(None, 36, 36, 128)	73856

```

-----
batch_normalization_2 (Batch Normalization) (None, 36, 36, 128) 512
-----
conv2d_3 (Conv2D) (None, 36, 36, 8) 1032
-----
batch_normalization_3 (Batch Normalization) (None, 36, 36, 8) 32
-----
flatten (Flatten) (None, 10368) 0
-----
dense (Dense) (None, 100) 1036900
-----
batch_normalization_4 (Batch Normalization) (None, 100) 400
-----
dense_1 (Dense) (None, 5) 505
-----
multivariate_normal_tri_l (Multivariate Normal Triangular) (M ((None, 2), (None, 2))) 0
-----
kl_divergence_add_loss (KLDivergence) (None, 2) 4
=====
Total params: 1,133,017
Trainable params: 1,132,353
Non-trainable params: 664
-----

```

```

In [18]: # decoder part
# inspired by https://www.tensorflow.org/tutorials/generative/cvae
decoder = Sequential([
    InputLayer(input_shape=(latent_dim,)),

    Dense(8*8*8),
    Reshape(target_shape=(8,8,8)),

    Conv2DTranspose(filters=32, kernel_size=(3,3), strides=(2,2), activation='relu'),
    BatchNormalization(),
    Conv2DTranspose(filters=16, kernel_size=(3,3), strides=(2,2), activation='relu'),
    BatchNormalization(),
    Conv2DTranspose(filters=1, kernel_size=(3, 3), strides=(2,2), activation='relu'),
    BatchNormalization(),

    Flatten(),
    Dense(tfpl.IndependentBernoulli.params_size(image_dim), activation='sigmoid'),
    tfpl.IndependentBernoulli(event_shape=image_dim)
], name='decoder')
# Print the model summary for the encoder and decoder networks.
decoder.summary()

Model: "decoder"
-----

```

Layer (type)	Output Shape	Param #
dense_2 (Dense)	(None, 512)	1536
reshape (Reshape)	(None, 8, 8, 8)	0
conv2d_transpose (Conv2DTran	(None, 16, 16, 32)	2336
batch_normalization_5 (Batch	(None, 16, 16, 32)	128
conv2d_transpose_1 (Conv2DTr	(None, 32, 32, 16)	4624
batch_normalization_6 (Batch	(None, 32, 32, 16)	64
conv2d_transpose_2 (Conv2DTr	(None, 64, 64, 1)	145
batch_normalization_7 (Batch	(None, 64, 64, 1)	4
flatten_1 (Flatten)	(None, 4096)	0
dense_3 (Dense)	(None, 3888)	15929136
independent_bernoulli (Indep	((None, 36, 36, 3), (None 0	
Total params: 15,937,973		
Trainable params: 15,937,875		
Non-trainable params: 98		

1.5 5. Train the variational autoencoder

- You should now train the variational autoencoder. Build the VAE using the `Model` class and the encoder and decoder models. Print the model summary.
- Compile the VAE with the negative log likelihood loss and train with the `fit` method, using the training and validation Datasets.
- Plot the learning curves for loss vs epoch for both training and validation sets.

```
In [19]: # Build the VAE using the Model class and the encoder and decoder models. Print the m
vae = Model(inputs=encoder.inputs, outputs=decoder(encoder.outputs), name='vae')
```

```
In [20]: # Compile the VAE with the negative log likelihood loss
def nll(y_true, y_pred):
    return -tf.reduce_mean(y_pred.log_prob(y_true))

vae.compile(loss=nll, optimizer='adam')
vae.summary()
```

Model: "vae"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 36, 36, 3)]	0
conv2d (Conv2D)	(None, 36, 36, 32)	896
batch_normalization (Batch Normalization)	(None, 36, 36, 32)	128
conv2d_1 (Conv2D)	(None, 36, 36, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 36, 36, 64)	256
conv2d_2 (Conv2D)	(None, 36, 36, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 36, 36, 128)	512
conv2d_3 (Conv2D)	(None, 36, 36, 8)	1032
batch_normalization_3 (Batch Normalization)	(None, 36, 36, 8)	32
flatten (Flatten)	(None, 10368)	0
dense (Dense)	(None, 100)	1036900
batch_normalization_4 (Batch Normalization)	(None, 100)	400
dense_1 (Dense)	(None, 5)	505
multivariate_normal_tri_l (Multivariate Normal Triangular Lower)	((None, 2), (None, 2))	0
kl_divergence_add_loss (KLDivergence)	(None, 2)	4
decoder (Sequential)	(None, 36, 36, 3)	15937973
Total params: 17,070,990		
Trainable params: 17,070,228		
Non-trainable params: 762		

```
In [21]: # and train with the fit method, using the training and validation Datasets.
         early_stopping = tf.keras.callbacks.EarlyStopping(monitor="val_loss",
                                                             min_delta=0.1, patience=5,
                                                             restore_best_weights=True)
```

```
history = vae.fit(train_set,
                  validation_data=val_set,
                  epochs=10,
                  callbacks=[early_stopping],
                  verbose=2)
```

Train for 37 steps, validate for 12 steps

Epoch 1/10

37/37 - 100s - loss: 263.2699 - val_loss: 1208.6427

Epoch 2/10

37/37 - 99s - loss: 92.1871 - val_loss: 590.4367

Epoch 3/10

37/37 - 98s - loss: 86.4432 - val_loss: 339.0957

Epoch 4/10

37/37 - 99s - loss: 86.7019 - val_loss: 250.3676

Epoch 5/10

37/37 - 96s - loss: 80.8608 - val_loss: 211.6612

Epoch 6/10

37/37 - 94s - loss: 85.6911 - val_loss: 167.8786

Epoch 7/10

37/37 - 94s - loss: 81.7220 - val_loss: 165.1337

Epoch 8/10

37/37 - 94s - loss: 82.0641 - val_loss: 153.0014


Epoch 9/10

37/37 - 94s - loss: 79.9850 - val_loss: 166.9210

Epoch 10/10

37/37 - 95s - loss: 79.4047 - val_loss: 150.0522

```
In [22]: # Plot the learning curves for loss vs epoch for both training and validation sets.
plt.plot(history.history["loss"])
plt.plot(history.history["val_loss"])
plt.show()
```



1.6 6. Use the encoder and decoder networks

- You can now put your encoder and decoder networks into practice!
- Randomly sample 1000 images from the dataset, and pass them through the encoder. Display the embeddings in a scatter plot (project to 2 dimensions if the latent space has dimension higher than two).
- Randomly sample 4 images from the dataset and for each image, display the original and reconstructed image from the VAE in a figure.
 - Use the mean of the output distribution to display the images.
- Randomly sample 6 latent variable realisations from the prior distribution, and display the images in a figure.
 - Again use the mean of the output distribution to display the images.


In [23]: *# Randomly sample 1000 images from the dataset, and pass them through the encoder.
Display the embeddings in a scatter plot (project to 2 dimensions if the latent spa*

```
idx = np.random.choice(np.arange(images.shape[0]), 1000)
embeddings = encoder(images[idx]/255.0).mean()
plt.scatter(embeddings[:,0], embeddings[:,1])
plt.show()
```

WARNING:tensorflow:Layer conv2d is casting an input tensor from dtype float64 to the layer's dtype

If you intended to run this layer in float32, you can safely ignore this warning. If in doubt,

To change all layers to have dtype float64 by default, call ``tf.keras.backend.set_floatx('float64')`



In [24]: *# Randomly sample 4 images from the dataset and for each image, display the original image and its reconstruction. Use the mean of the output distribution to display the images.*

```
idx = np.random.choice(np.arange(images.shape[0]), 4)
reconstructions = vae(images[idx]).mean().numpy()

plt.figure(figsize=(15, 6))
for i in range(4):
    plt.subplot(2, 4, i+1)
    plt.imshow(images[idx[i]])
    plt.title("Original: {}".format(i+1))
    plt.axis("off")

    plt.subplot(2, 4, i+5)
    plt.imshow(reconstructions[i])
    plt.title("Reconstruction: {}".format(i+1))
    plt.axis("off")
plt.show()
```

```

_FallbackException                                Traceback (most recent call last)

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/gen_nn_ops.py in con
925         "explicit_paddings", explicit_paddings, "data_format", data_format,
--> 926         "dilations", dilations)
927     return _result

```

_FallbackException: Expecting int64_t value for attr strides, got numpy.int32

During handling of the above exception, another exception occurred:

```

ValueError                                         Traceback (most recent call last)

<ipython-input-24-9a47393497c4> in <module>
      3
      4 idx = np.random.choice(np.arange(images.shape[0]), 4)
----> 5 reconstructions = vae(images[idx]).mean().numpy()
      6
      7 plt.figure(figsize=(15, 6))

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/keras/engine/base_layer.py
820         with base_layer_utils.autocast_context_manager(
821             self._compute_dtype):
--> 822             outputs = self.call(cast_inputs, *args, **kwargs)
823             self._handle_activity_regularization(inputs, outputs)
824             self._set_mask_metadata(inputs, outputs, input_masks)

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/keras/engine/network.py :
715     return self._run_internal_graph(
716         inputs, training=training, mask=mask,
--> 717         convert_kwargs_to_constants=base_layer_utils.call_context().saving)
718
719     def compute_output_shape(self, input_shape):

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/keras/engine/network.py :
889
890     # Compute outputs.
--> 891     output_tensors = layer(computed_tensors, **kwargs)
892
893     # Update tensor_dict.

```

```

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/keras/engine/base_layer.py
820         with base_layer_utils.autocast_context_manager(
821             self._compute_dtype):
--> 822             outputs = self.call(cast_inputs, *args, **kwargs)
823             self._handle_activity_regularization(inputs, outputs)
824             self._set_mask_metadata(inputs, outputs, input_masks)

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/keras/layers/convolutional.py
207         inputs = array_ops.pad(inputs, self._compute_causal_padding())
208
--> 209         outputs = self._convolution_op(inputs, self.kernel)
210
211         if self.use_bias:

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/nn_ops.py in __call__(self, inp, filter, call_from_convolution=False)
1133         call_from_convolution=False)
1134     else:
-> 1135         return self.conv_op(inp, filter)
1136     # copybara:strip_end
1137     # copybara:insert return self.conv_op(inp, filter)

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/nn_ops.py in __call__(self, inp, filter)
638
639     def __call__(self, inp, filter): # pylint: disable=redefined-builtin
--> 640         return self.call(inp, filter)
641
642

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/nn_ops.py in __call__(self, inp, filter, padding=self.padding, data_format=self.data_format, name=self.name)
237         padding=self.padding,
238         data_format=self.data_format,
--> 239         name=self.name)
240
241

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/nn_ops.py in conv2d(self, inp, filter, data_format=data_format, dilations=dilations, name=name)
2009         data_format=data_format,
2010         dilations=dilations,
-> 2011         name=name)
2012
2013

```

```

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/gen_nn_ops.py in con
931         input, filter, strides=strides, use_cudnn_on_gpu=use_cudnn_on_gpu,
932         padding=padding, explicit_paddings=explicit_paddings,
--> 933         data_format=data_format, dilations=dilations, name=name, ctx=_ctx)
934     except _core._SymbolicException:
935         pass # Add nodes to the TensorFlow graph.

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/gen_nn_ops.py in con
1013         "'conv2d' Op, not %r." % dilations)
1014     dilations = [_execute.make_int(_i, "dilations") for _i in dilations]
-> 1015     _attr_T, _inputs_T = _execute.args_to_matching_eager([input, filter], ctx)
1016     (input, filter) = _inputs_T
1017     _inputs_flat = [input, filter]

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/eager/execute.py in args
265         dtype = ret[-1].dtype
266     else:
--> 267         ret = [ops.convert_to_tensor(t, dtype, ctx=ctx) for t in l]
268
269     # TODO(slebedev): consider removing this as it leaks a Keras concept.

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/eager/execute.py in <list
265         dtype = ret[-1].dtype
266     else:
--> 267         ret = [ops.convert_to_tensor(t, dtype, ctx=ctx) for t in l]
268
269     # TODO(slebedev): consider removing this as it leaks a Keras concept.

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/framework/ops.py in conv
1312
1313     if ret is None:
-> 1314         ret = conversion_func(value, dtype=dtype, name=name, as_ref=as_ref)
1315
1316     if ret is NotImplemented:

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/resource_variable_ops
1792
1793 def _dense_var_to_tensor(var, dtype=None, name=None, as_ref=False):
-> 1794     return var._dense_var_to_tensor(dtype=dtype, name=name, as_ref=as_ref) # pylint
1795
1796

```

```

/opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/resource_variable_ops.py:1213: raise ValueError(
1214     "Incompatible type conversion requested to type {!r} for variable "
-> 1215     "of type {!r}".format(dtype.name, self.dtype.name))
1216     if as_ref:
1217         return self.read_value().op.inputs[0]

```

ValueError: Incompatible type conversion requested to type 'uint8' for variable of type

In [25]: *# Randomly sample 6 latent variable realisations from the prior distribution, and display them. Again use the mean of the output distribution to display the images.*

```

latent_variables = np.random.uniform(-2, 2, (6, latent_dim))
realisations = decoder(latent_variables).mean()

plt.figure(figsize=(15, 6))
for i in range(6):
    plt.subplot(1, 6, i+1)
    plt.imshow(realisations[i])
    plt.title("Realisation: {}".format(i+1))
    plt.axis("off")
plt.show()

```

WARNING:tensorflow:Layer dense_2 is casting an input tensor from dtype float64 to the layer's dtype float32. This operation can have a negative effect on performance.

If you intended to run this layer in float32, you can safely ignore this warning. If in doubt, please open an issue at <https://github.com/tensorflow/tensorflow>.

To change all layers to have dtype float64 by default, call `tf.keras.backend.set_floatx('float64')`.



1.7 Make a video of latent space interpolation (not assessed)

- Just for fun, you can run the code below to create a video of your decoder's generations, depending on the latent space.

In [26]: *# Function to create animation*

```
import matplotlib.animation as anim
from IPython.display import HTML

def get_animation(latent_size, decoder, interpolation_length=500):
    assert latent_size >= 2, "Latent space must be at least 2-dimensional for plotting"
    fig = plt.figure(figsize=(9, 4))
    ax1 = fig.add_subplot(1,2,1)
    ax1.set_xlim([-3, 3])
    ax1.set_ylim([-3, 3])
    ax1.set_title("Latent space")
    ax1.axes.get_xaxis().set_visible(False)
    ax1.axes.get_yaxis().set_visible(False)
    ax2 = fig.add_subplot(1,2,2)
    ax2.set_title("Data space")
    ax2.axes.get_xaxis().set_visible(False)
    ax2.axes.get_yaxis().set_visible(False)

    # initializing a line variable
    line, = ax1.plot([], [], marker='o')
    img2 = ax2.imshow(np.zeros((36, 36, 3)))

    freqs = np.random.uniform(low=0.1, high=0.2, size=(latent_size,))
    phases = np.random.randn(latent_size)
    input_points = np.arange(interpolation_length)
    latent_coords = []
    for i in range(latent_size):
        latent_coords.append(2 * np.sin((freqs[i]*input_points + phases[i])).astype(np.float32))

    def animate(i):
        z = tf.constant([coord[i] for coord in latent_coords])
        img_out = np.squeeze(decoder(z[np.newaxis, ...]).mean().numpy())
        line.set_data(z.numpy()[0], z.numpy()[1])
        img2.set_data(np.clip(img_out, 0, 1))
        return (line, img2)

    return anim.FuncAnimation(fig, animate, frames=interpolation_length,
                              repeat=False, blit=True, interval=150)
```

In [27]: *# Create the animation*

```
a = get_animation(latent_size, decoder, interpolation_length=200)
HTML(a.to_html5_video())
```

NameError

Traceback (most recent call last)

```
<ipython-input-27-f8780fad386d> in <module>
      1 # Create the animation
      2
----> 3 a = get_animation(latent_size, decoder, interpolation_length=200)
      4 HTML(a.to_html5_video())
```

NameError: name 'latent_size' is not defined

In []: